

## Programming 2 Report

### Introduction

I chose to do the Optical Flow exercise. I'm turning in one python code named "OpticalFlow.py", which contains the Lucas-Kanade implementation. We were given frame1\_a.png, frame1\_b.png, frame2\_a.png and frame2\_b.png as our dataset. I have hard coded the image paths to be the same location as the code file. Running the code outputs the results for optical flow on top of frame1\_b.png and frame2\_b.png, with relevant  $V_x$  and  $V_y$  vectors shown in magenta and some of the magnitude as green arrows on the image. For the write-up, I commented out the portion I did not need and the magnitude arrows were scaled so they are more visible. I also have a function called graph() that graphs the  $V_x$  and  $V_y$  vectors (as  $V$ ) only and has arrow heads that show the direction of  $V$  for each pixel point.

Each image is preprocessed by adding zero pads of width 1 and the intensity (the pixel value) of each pixel is normalized by dividing by 255. The gradient filters  $g_x$  and  $g_y$  as given in the assignment sheet were used to calculate  $I_x$  and  $I_y$ . It was calculated as  $I_t$  is calculated by subtracting frame1's intensity from frame 2. Using the given equation in the assignment sheet,  $V_x$  and  $V_y$  were calculated. My implementation of Lucas-Kanade is shown below:


```
OpticalFlow4.py
def lucas_kanade(V, padded1, padded2): # Lucas-Kanade Function
    r_count = 0
    for row in padded1:
        end = len(row)
        if np.array_equal(row, padded1[0]): # zero padded row
            pass # Do nothing
        else:
            stop = end-1
            col = 1
            while col < stop:
                window1 = np.array([[padded1[r_count-1][col-1], padded1[r_count-1][col], padded1[r_count-1][col+1]],
                                    [row[col-1], row[col], row[col+1]],
                                    [padded1[r_count+1][col-1], padded1[r_count+1][col], padded1[r_count+1][col+1]])
                window2 = np.array([[padded2[r_count-1][col-1], padded2[r_count-1][col], padded2[r_count-1][col+1]],
                                    [padded2[r_count][col-1], padded2[r_count][col], padded2[r_count][col+1]],
                                    [padded2[r_count+1][col-1], padded2[r_count+1][col], padded2[r_count+1][col+1]])
                w1 = 1/255 * window1 # Normalization
                w2 = 1/255 * window2 # Normalization
                I_x = w1 * DoG_gx # Element-wise multiplication
                I_y = w1 * DoG_gy # Element-wise multiplication
                I_t = np.subtract(w2, w1) # Element-wise subtraction
                A = np.array([[np.sum(I_x * I_x), np.sum(I_x * I_y)],
                              [np.sum(I_x * I_y), np.sum(I_y * I_y)]])
                A_inv = np.linalg.pinv(A) # A^(-1)
                b = np.zeros((2, 1))
                b[0][0] = -1 * np.sum(I_x * I_t)
                b[1][0] = -1 * np.sum(I_y * I_t)
                u_v = np.matmul(A_inv, b)
                V[r_count-1][col-1][0] = u_v[0]
                V[r_count-1][col-1][1] = u_v[1]
                col = col + 1
            r_count = r_count + 1
```

## Results


### a) Image 1

Below are the original images of the two frames. Frame 1 is frame1\_a.png and Frame 2 is frame1\_b.png. We can tell with our eyes that the ball and the hands and arms of the two people move.

#### Image 1 Frame 1:



#### Image 1 Frame 2:



**V<sub>x</sub> and V<sub>y</sub> values for Image 1:**

```
Image 1 Vx and Vy: [[[-2.02765512e-02  4.21589235e-03]
 [ 2.41784896e-02 -1.79814989e-02]
 [ 3.35874974e-02 -4.43205907e-02]
 ...
 [-2.37262020e-03  1.84614934e-05]
 [-2.35189191e-03  5.36301255e-05]
 [ 1.28403050e-02 -1.81126726e-03]]

 [[-6.89679548e-02  1.48259947e-02]
 [ 2.66109639e-02  1.55784158e-02]
 [ 4.64741833e-02  5.11488020e-03]
 ...
 [-5.83672342e-03  5.97351090e-03]
 [ 9.19409882e-04  6.78734901e-03]
 [ 9.36752803e-03  7.80227483e-03]]

 [[-3.91630707e-02  4.75342576e-02]
 [ 6.42770283e-03  5.14184975e-02]
 [ 2.06727587e-02  2.57155237e-02]
 ...
 [-2.05219806e-03  2.11329519e-03]
 [ 3.04297446e-03  2.94911584e-03]
 [-3.63847909e-03  2.20003333e-03]]

 ...


 [[ 0.00000000e+00  2.63825122e-03]
 [ 3.22810578e-03  1.12016947e-03]
 [-3.28953720e-03  3.26037997e-03]
 ...
 [-4.45259298e-03  4.45259298e-03]
 [ 2.36541599e-02 -4.45259298e-03]
 [ 0.00000000e+00 -1.14864865e-02]]

 [[-1.14463298e-05  2.60415224e-03]
 [ 2.18796759e-03  2.17855617e-03]
 [-6.57857102e-03  2.17378868e-03]
 ...
 [-4.26386551e-03  4.80763930e-03]
 [ 1.44100054e-02 -4.62207722e-03]
 [ 9.25925926e-03 -1.75675676e-02]]

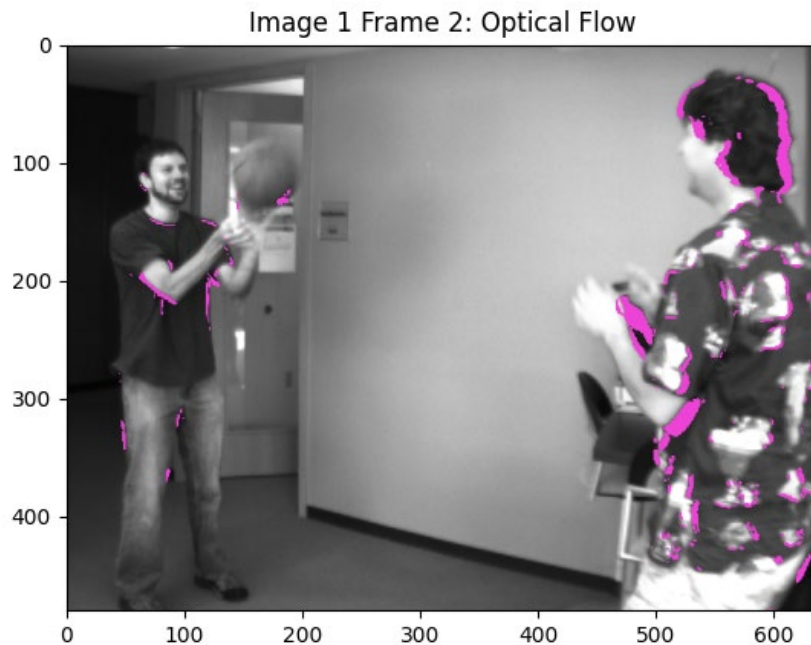
 [[ 6.56464571e-03  6.50743909e-03]
 [-6.52635839e-03  2.20239117e-03]
 [-7.84861151e-03  4.31380149e-03]
 ...
 [-5.40254237e-03 -9.00423729e-03]
 [ 0.00000000e+00  0.00000000e+00]
 [ 2.32397462e-02 -5.08762012e-03]]]
```

### Image 1 Optical Flow $V_x$ and $V_y$ plot on image:

My values for  $V_x$  and  $V_y$  were mostly very small, and hard to visualize. I plotted  $V_x$  and  $V_y$  as magenta lines from pixel  $(x, y)$  using OpenCV's `arrowedLine()` function. My Lucas-Kanade Implementation detected a lot of noise, so I tried to limit the number of points plotted. Below is after ignoring  $V_x$  and  $V_y$  values between  $-0.1$  and  $0.1$ :




Below is after ignoring  $V_x$  and  $V_y$  values between  $-0.2$  and  $0.2$ :




Below is after ignoring  $V_x$  and  $V_y$  values between -0.3 and 0.3:




Below is after ignoring  $V_x$  and  $V_y$  values between -0.5 and 0.5




I have also included a `graph()` function that uses `matplotlib` that is able to plot the  $V_x$  and  $V_y$  (as  $V$ ) in a coordinate. The result, without any scaling was below:



Because it's still hard to tell the direction, I plotted every 10<sup>th</sup> pixel point and scaled the arrows to be 10 times longer for below result:




Ignoring  $V_x$  and  $V_y$  between -0.3 and 0.3 to get rid of some noise gives results below:





### Image 1 Optical Flow Magnitude:

I decided ignoring  $V_x$  and  $V_y$  between -0.3 and 0.3 looked good and used it for drawing magnitudes (green arrows). To make the arrows more visible, I multiplied the  $V_x$  and  $V_y$  values by 150 and ignored magnitudes less than 0.2. I also skipped drawing arrows on every 3<sup>rd</sup> row and every 4<sup>th</sup> column in the image.







## b) Image 2

Below are the original images of the two frames for image 2. Frame 1 is frame2\_a.png and Frame 2 is frame2\_b.png. This one is hard to tell with eyes how the leaves move.

### Image 2 Frame 1:



### Image 2 Frame 2:



**V<sub>x</sub> and V<sub>y</sub> values for Image 2:**

```
Image 2 Vx and Vy: [[[-7.16741731e-03  1.45327234e-03]
 [ 2.56773928e-03 -1.42776310e-03]
 [ 4.20775615e-03 -1.41242938e-03]
 ...
 [-1.70302297e-01 -6.96202899e-02]
 [-3.76550243e-01  1.44141815e-01]
 [ 3.91362924e-01 -1.88455710e-01]]]

 [[-5.65277988e-03  3.44090095e-03]
 [ 2.81831387e-03  2.79783792e-03]
 [ 2.07286532e-03  6.76219370e-04]
 ...
 [-3.04052846e-01 -1.27973637e-01]
 [-3.12992261e-01  2.94818124e-01]
 [ 4.50748995e-01  2.77522894e-01]]]

 [[-4.21341665e-03  2.54519161e-03]
 [ 1.38855487e-03  2.80883244e-03]
 [-3.10795770e-05  1.39658483e-03]
 ...
 [-2.02535138e-01  2.35705817e-01]
 [-1.92080115e-01  3.22056947e-01]
 [ 4.09892051e-01  2.11977626e-01]]]

 ...

 [[-3.47893388e-01  3.37896835e-01]
 [ 1.10344104e-01  2.97924701e-01]
 [ 9.52202828e-02  2.87161226e-01]
 ...
 [ 1.20733466e-02  8.88656948e-03]
 [ 7.87593112e-03  1.10060654e-03]
 [-1.25758711e-02  1.59840086e-02]]]


 [[-3.46733223e-01  2.92509916e-01]
 [ 9.52929381e-02  2.81664292e-01]
 [ 7.95735215e-02  2.77271923e-01]
 ...
 [ 2.06762837e-02  1.01314636e-02]
 [ 2.93636286e-02  8.84782609e-03]
 [-1.97309960e-02  7.71696517e-03]]]

 [[-2.66151568e-01  3.51518359e-01]
 [ 7.54871384e-03  4.46770981e-01]
 [ 4.04632523e-03  4.45776225e-01]
 ...
 [ 1.41502289e-02  1.57862366e-02]
 [ 2.92322507e-02 -2.59269795e-02]
 [ 4.17055344e-03 -5.76948791e-02]]]
```

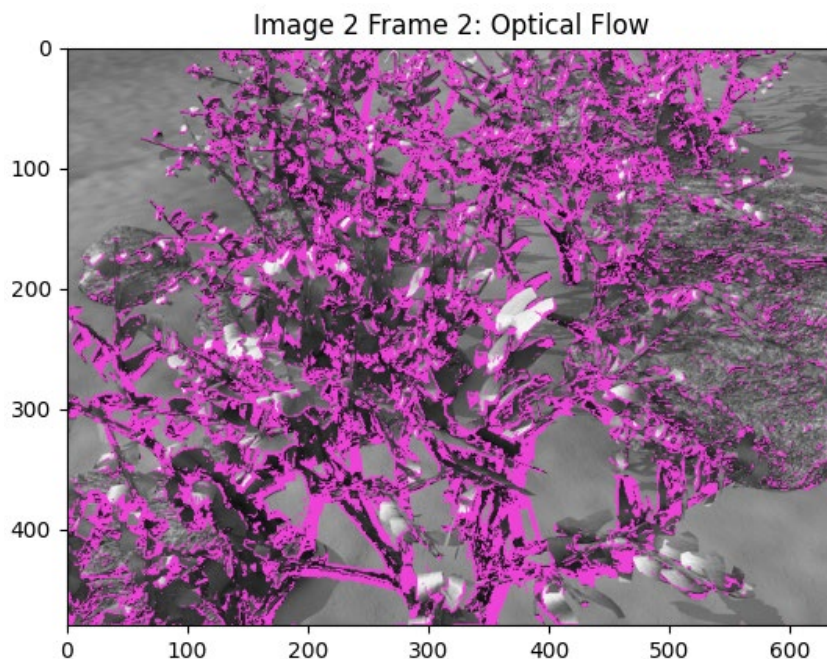
□

### Image 2 Optical Flow $V_x$ and $V_y$ plot:

I plotted  $V_x$  and  $V_y$  as magenta lines from pixel  $(x, y)$ . My Lucas-Kanade Implementation detected a lot of noise, so I tried to limit the number of points plotted. Below is after ignoring  $V_x$  and  $V_y$  values between  $-0.1$  and  $0.1$ :

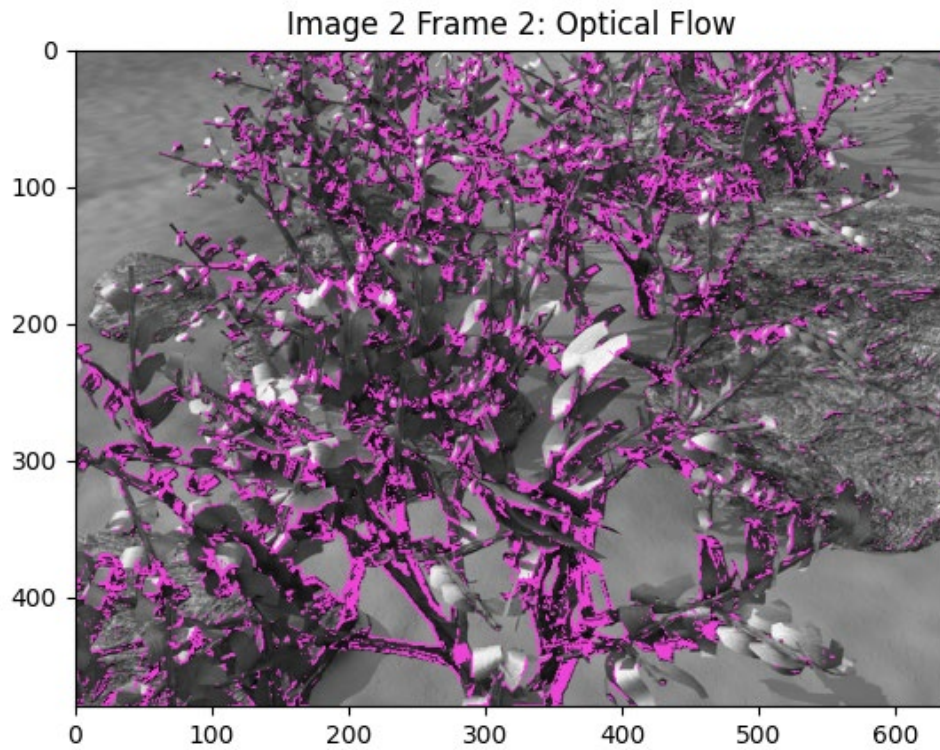


Below is after ignoring  $V_x$  and  $V_y$  values between  $-0.2$  and  $0.2$ :







Below is after ignoring  $V_x$  and  $V_y$  values between -0.3 and 0.3:




Below is after ignoring  $V_x$  and  $V_y$  values between -0.5 and 0.5:




The result from the `graph()` function, without any scaling was below:



Because it's still hard to tell the direction, I plotted every 10<sup>th</sup> pixel point and scaled the arrows to be 10 times longer for below result:




Ignoring  $V_x$  and  $V_y$  between -0.5 and 0.5 to get rid of some noise gives results below:






## Image 2 Optical Flow Magnitude:

I decided ignoring  $V_x$  and  $V_y$  between  $-0.5$  and  $0.5$  looked good and used it for drawing magnitudes (green arrows) for image 2. To make the arrows more visible, I multiplied the  $V_x$  and  $V_y$  values by 4 and ignored magnitudes less than 1. I also skipped drawing arrows on every 3<sup>rd</sup> row and every 4<sup>th</sup> column in the image.



Below is another result using different scaling and plotting limit as an example:



The general movement direction is the same, even though different points' magnitudes are shown.

### **Conclusion**

My implementation of Lucas-Kanade Optical Flow method can find the objects that move and ignore the background. It picked up a lot of noise, and perhaps I should have rounded off some of the values in the matrices. Ignoring some of the  $V_x$  and  $V_y$  values that were closer to zero helped dealing with noise. I think the general direction of the movements graphed seems correct as most movements are seen around the persons' arms in Image 1, and the person on the right is moving to the right side of the frame. For image 2, the algorithm focused mostly on the tree and its shadow, which are what you would expect to move.