

Programming 5 Report

Introduction

I worked on exercise 2 and turning in two python codes named “GAN.py” and “ConvGAN.py”. The dataset used is the Fashion-MNIST dataset. The “GAN.py” code outputs the summary for the generator, the discriminator, and the combined GAN model for step 1, trains the models by first training the discriminator with unfrozen weights for an epoch and then training the GAN model’s generator with the discriminator’s weights frozen for an epoch. Both are trained 100 epochs each. The code also saves the first 3 images from the GAN model’s generator per 10 epochs and the last epoch 99. The “ConvGAN.py” code outputs the summary for the CNN generator, the CNN discriminator, and the combined CNN GAN model for step 3, trains the models by first training the discriminator with unfrozen weights for an epoch and then training the GAN model’s generator with the discriminator’s weights frozen for an epoch. Both are trained 100 epochs each. The code saves the first 3 images from the step 3 GAN model’s generator per 10 epochs and the last epoch 99.

Results

Step 1: Design the GAN

I experimented with my GAN model’s layer orders, the number of dense layers, width of the dense layers, and settled on what’s shown below, though my GAN model does not seem to produce a fake image that looks like it could be one of the real images, it does come close. I also experimented with batch size but didn’t see too much improvement, so I stuck with 100. I think my generator is weaker than the discriminator, but I could not figure out how to balance them well. I used batch normalization between dense layers and leaky RELU as suggested.

The Generator

```
# Generator
gen_model = models.Sequential()
gen_model.add(tf.keras.Input(shape=(100,)))
gen_model.add(layers.Dense(256))
gen_model.add(layers.LeakyReLU())
gen_model.add(layers.BatchNormalization())
gen_model.add(layers.Dense(512))
gen_model.add(layers.LeakyReLU())
gen_model.add(layers.BatchNormalization())
gen_model.add(layers.Dense(1024))
gen_model.add(layers.LeakyReLU())
gen_model.add(layers.BatchNormalization())
gen_model.add(layers.Dense(784)) # (None, 784)
gen_model.add(layers.LeakyReLU()) # (None, 28, 28, 1)
gen_model.add(layers.Reshape((28, 28, 1), input_shape=(2, )))
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 256)	25856
leaky_re_lu (LeakyReLU)	(None, 256)	0
batch_normalization (Batch Normalization)	(None, 256)	1024
dense_1 (Dense)	(None, 512)	131584
leaky_re_lu_1 (LeakyReLU)	(None, 512)	0
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dense_2 (Dense)	(None, 1024)	525312
leaky_re_lu_2 (LeakyReLU)	(None, 1024)	0
batch_normalization_2 (Batch Normalization)	(None, 1024)	4096
dense_3 (Dense)	(None, 784)	803600
leaky_re_lu_3 (LeakyReLU)	(None, 784)	0
reshape (Reshape)	(None, 28, 28, 1)	0

```

=====
Total params: 1,493,520
Trainable params: 1,489,936
Non-trainable params: 3,584

```

The Discriminator

```

1 # Discriminator
2 disc_model = models.Sequential()
3 disc_model.add(tf.keras.Input(shape=(28, 28, 1)))
4 disc_model.add(layers.Flatten()) # (None, 784)
5 disc_model.add(layers.Dense(1024))
6 disc_model.add(layers.LeakyReLU())
7 disc_model.add(layers.BatchNormalization())
8 disc_model.add(layers.Dense(512))
9 disc_model.add(layers.LeakyReLU())
10 disc_model.add(layers.BatchNormalization())
11 disc_model.add(layers.Dense(256))
12 disc_model.add(layers.LeakyReLU())
13 disc_model.add(layers.BatchNormalization())
14 disc_model.add(layers.Dense(1, activation='sigmoid')) # (None, 1)

```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense_4 (Dense)	(None, 1024)	803840
leaky_re_lu_4 (LeakyReLU)	(None, 1024)	0
batch_normalization_3 (Batch Normalization)	(None, 1024)	4096
dense_5 (Dense)	(None, 512)	524800
leaky_re_lu_5 (LeakyReLU)	(None, 512)	0
batch_normalization_4 (Batch Normalization)	(None, 512)	2048
dense_6 (Dense)	(None, 256)	131328
leaky_re_lu_6 (LeakyReLU)	(None, 256)	0
batch_normalization_5 (Batch Normalization)	(None, 256)	1024
dense_7 (Dense)	(None, 1)	257

```
=====  
Total params: 1,467,393  
Trainable params: 1,463,809  
Non-trainable params: 3,584
```

The GAN Model

```
# Full GAN  
disc_model.trainable = False  
GAN_model = models.Sequential([tf.keras.layers.Input(  
    shape=(100,)), gen_model, disc_model])  
GAN_model.build(input_shape=(100,))  
GAN_model.summary()
```

```
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
sequential (Sequential)	(None, 28, 28, 1)	1493520
sequential_1 (Sequential)	(None, 1)	1467393

```
=====  
Total params: 2,960,913  
Trainable params: 1,489,936  
Non-trainable params: 1,470,977
```

Step 2: Training GAN model and Discriminator model

Data was preprocessed by multiplying the image values by 1/255. I trained my model for 100 epochs with 100 batch size. Optimizer was adam and the loss measured was binary cross-

entropy. I coded the training loop following the guide in

https://keras.io/guides/writing_a_training_loop_from_scratch/

Training Loss per Epoch


Epoch	Discriminator Training Loss	GAN Training Loss	Epoch	Discriminator Training Loss	GAN Training Loss
0	3.107323565121334e-10	1.8828283487344882e-20	50	0.0584297813475132	0.0009730348247103393
1	6.628647042816738e-06	0.0	51	0.0027646152302622795	0.0021094121038913727
2	1.7890829334503422e-13	0.0	52	0.0011166115291416645	0.0005461795371957123
3	1.2568125779005338e-25	0.0	53	0.002532252110540867	9.268229769077152e-05
4	1.181110646396125e-19	0.0	54	0.0027956985868513584	0.00019525884999893606
5	1.5007686689514291e-21	0.0	55	0.00580389192327857	0.0002568997151684016
6	1.313190178819923e-08	0.0	56	0.0020593523513525724	1.0518755516386591e-05
7	6.885981774730923e-38	0.0	57	0.0004420839250087738	0.00017423676035832614
8	3.319027896964144e-08	0.0	58	9.997850611398462e-06	0.00038750431849621236
9	0.00014608855417463928	0.0	59	4.398263627081178e-06	2.378181918061273e-08
10	1.76285345787619e-06	1.3902207612991333	60	9.690074512036517e-05	8.186914055841044e-05
11	2.6132571292691864e-05	6.445011635491937e-09	61	4.909489143756218e-05	2.6243464162689634e-05
12	6.386033055605367e-05	0.0001176200748886913	62	0.0003471345407888293	0.0007622029515914619
13	0.00033392079058103263	0.3596891164779663	63	0.0012816556263715029	0.000492788094561547
14	0.0013279863633215427	0.2481156885623932	64	0.0030268938280642033	0.0005286442465148866
15	0.019973058253526688	0.040013059973716736	65	0.047363437712192535	3.6267261748434976e-05
16	0.05885313078761101	0.4276794195175171	66	1.629531107028015e-05	5.95403544139117e-05
17	0.01747729629278183	0.1900664120912552	67	1.4044260296941502e-06	3.817094210156746e-12
18	0.0020217604469507933	0.047411952167749405	68	1.0396170324611376e-07	4.421699117962419e-11
19	0.009358644485473633	0.14184927940368652	69	2.8612439564312808e-05	2.594969373603817e-06
20	0.0202618557959795	0.0028923461213707924	70	3.4176098324678605e-06	8.584412825030086e-08



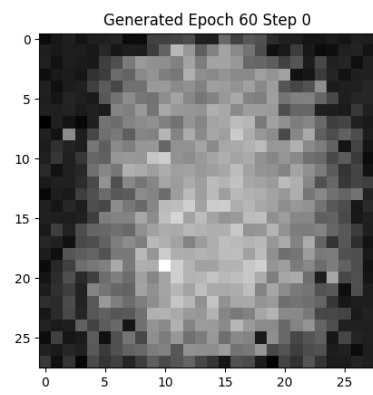








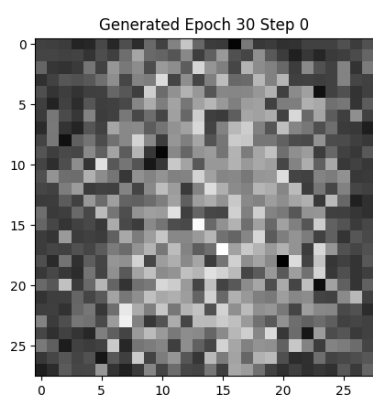
21	0.023538520559 66854	0.0349653102457 52335	71	0.000131845881696 79046	3.396354486540076 e-08
22	0.025270387530 326843	0.0219049230217 93365	72	7.323890258703614 e-06	1.742265158100053 7e-05
23	0.007923358120 024204	0.1110096946358 6807	73	0.003940358292311 43	0.001423598267138 0043
24	0.055959206074 47624	0.2146822959184 6466	74	0.000147764396388 0837	0.001980350119993 0906
25	0.025974364951 252937	0.0041482583619 65418	75	0.003828159766271 7104	6.568321259692311 e-05
26	0.012430590577 423573	0.0073649422265 58924	76	0.039498951286077 5	0.010636317543685 436
27	0.019361548125 743866	0.0092046102508 90255	77	0.002626393921673 298	0.005469164811074 734
28	0.010745820589 363575	0.1221443712711 3342	78	0.015969334170222 282	0.001841038465499 878
29	0.022217316552 996635	0.0450667664408 6838	79	0.045453798025846 48	0.006780379917472 601
30	0.010732552967 965603	0.0183046981692 31415	80	0.004218677990138 531	0.016124464571475 983
31	0.037731196731 328964	0.0285971108824 01466	81	0.007464053574949 503	0.006814803462475 538
32	0.022078871726 989746	0.0196036845445 63293	82	0.016049819067120 552	0.007481928449124 098
33	0.037839341908 693314	0.0250635817646 9803	83	0.006777258589863 777	0.004113648552447 5574
34	0.018960848450 660706	0.0270377397537 23145	84	0.005471150856465 101	0.000305945053696 6324
35	0.018887551501 393318	0.0212625600397 58682	85	0.004278677515685 558	0.004491932690143 585
36	0.013619081117 212772	0.0092500541359 18617	86	0.011253374628722 668	0.002120351186022 1624
37	0.013600916601 717472	0.0172690153121 94824	87	0.006726287771016 359	0.011302080005407 333
38	0.013729634694 755077	0.0126755535602 56958	88	0.003075928660109 639	0.006646629422903 061
39	0.013807655312 120914	0.0019850856624 543667	89	0.030716111883521 08	0.001207215827889 7405
40	0.010182742960 751057	0.0136740412563 08556	90	0.003150765784084 797	0.006561307702213 526
41	0.007732328958 809376	0.0046906964853 40595	91	0.017119912430644 035	0.009423914365470 41
42	0.006054171826 690435	0.0059003615751 86253	92	0.003734018886461 854	0.002064297907054 4243
43	0.006444049999 117851	0.0023717170115 560293	93	0.006089173723012 209	0.001389518962241 7092
44	0.003491320880 1299334	0.0245840121060 60982	94	0.019408116117119 79	0.007762295193970 203
45	0.005997875705 361366	0.0091615654528 14102	95	0.002697514835745 096	0.004071711096912 6225


46	0.003262854181 230068	0.0026606740429 997444	96	0.009124013595283 031	0.008223935961723 328
47	0.002659046091 1393166	0.0012421151623 129845	97	0.021852925419807 434	0.008333140984177 59
48	0.007605062332 004309	0.0023565338924 52717	98	0.035771340131759 644	0.011003077030181 885
49	0.005776549689 471722	0.0004685789463 110268	99	0.016789272427558 9	0.003909039776772 261

Synthetic Images

I grabbed the first 3 synthetic images every 10th epoch right before training:







Step 3: Conv GAN

I experimented with layer depth and width as well for CNN GAN model's generator and discriminator. I used batch normalization and Leaky RELU. I tried RELU and it did not work well. I tried to do stride 2 and filter size 5 as suggested for all the layers but had a hard time making the generator produce a 28x28x1 output that way, so I changed to filter size 2 with stride 2. None of what I tried made my CNN GAN model as good as the GAN model from step 1 and the below is the best model I had from the experiments:

The CNN Generator

```
cnn_gen_model = models.Sequential()  
cnn_gen_model.add(tf.keras.Input(shape=(6272,)))  
cnn_gen_model.add(layers.Reshape(  
    (7, 7, 128), input_shape=(6272, ))) # 28 * 28 * 8  
cnn_gen_model.add(layers.Conv2DTranspose(56, 2, strides=2))  
cnn_gen_model.add(layers.LeakyReLU())  
cnn_gen_model.add(layers.BatchNormalization()) # (None, 14, 14, 56)  
cnn_gen_model.add(layers.Conv2DTranspose(112, 2, strides=2))  
cnn_gen_model.add(layers.LeakyReLU())  
cnn_gen_model.add(layers.BatchNormalization()) # (None, 28, 28, 112)  
cnn_gen_model.add(layers.Conv2D(1, 1)) # (None, 28, 28, 1)
```

Model: "sequential"

Layer (type)	Output Shape	Param #
reshape (Reshape)	(None, 7, 7, 128)	0
conv2d_transpose (Conv2DTranspose)	(None, 14, 14, 56)	28728
leaky_re_lu (LeakyReLU)	(None, 14, 14, 56)	0
batch_normalization (BatchNormalization)	(None, 14, 14, 56)	224
conv2d_transpose_1 (Conv2DTranspose)	(None, 28, 28, 112)	25200
leaky_re_lu_1 (LeakyReLU)	(None, 28, 28, 112)	0
batch_normalization_1 (BatchNormalization)	(None, 28, 28, 112)	448
conv2d (Conv2D)	(None, 28, 28, 1)	113
Total params: 54,713		
Trainable params: 54,377		
Non-trainable params: 336		

The CNN Discriminator

```
cnn_disc_model = models.Sequential()  
cnn_disc_model.add(tf.keras.Input(shape=(28, 28, 1)))  
cnn_disc_model.add(layers.Conv2D(32, 5, strides=2))  
cnn_disc_model.add(layers.LeakyReLU())  
cnn_disc_model.add(layers.BatchNormalization())  
cnn_disc_model.add(layers.Conv2D(64, 5, strides=2))  
cnn_disc_model.add(layers.LeakyReLU())  
cnn_disc_model.add(layers.BatchNormalization())  
cnn_disc_model.add(layers.Flatten())  
cnn_disc_model.add(layers.Dense(1, activation='sigmoid')) # (None, 1)
```

```

Model: "sequential_1"
-----
Layer (type)                 Output Shape                 Param #
-----
conv2d_1 (Conv2D)            (None, 12, 12, 32)          832
leaky_re_lu_2 (LeakyReLU)    (None, 12, 12, 32)          0
batch_normalization_2 (Batc  (None, 12, 12, 32)          128
hNormalization)
conv2d_2 (Conv2D)            (None, 4, 4, 64)            51264
leaky_re_lu_3 (LeakyReLU)    (None, 4, 4, 64)            0
batch_normalization_3 (Batc  (None, 4, 4, 64)            256
hNormalization)
flatten (Flatten)            (None, 1024)                 0
dense (Dense)                 (None, 1)                    1025
-----
Total params: 53,505
Trainable params: 53,313
Non-trainable params: 192

```

The CNN GAN Model

```

# Full GAN
cnn_disc_model.trainable = False
cnn_GAN_model = models.Sequential([tf.keras.layers.Input(
    shape=(6272,)), cnn_gen_model, cnn_disc_model])
cnn_GAN_model.build(input_shape=(6272,))

```

```

Model: "sequential_2"
-----
Layer (type)                 Output Shape                 Param #
-----
sequential (Sequential)      (None, 28, 28, 1)           54713
sequential_1 (Sequential)    (None, 1)                    53505
-----
Total params: 108,218
Trainable params: 54,377
Non-trainable params: 53,841

```


I trained my model for 100 epochs with batch size 100. Optimizer chosen was adam and the loss measured was binary cross-entropy.

CNN Training Loss Per Epoch

Epoch	Training Loss	Test Loss	Epoch	Training Loss	Test Loss
0	0.0001791047106962651	1.2611647005344524e-34	50	0.0005514330696314573	8.770377159118652
1	0.05957120284438133	0.0015854808734729886	51	0.001434166100807488	6.432137966156006
2	0.010034257546067238	4.5194802567039005e-08	52	4.7547848225804046e-05	5.066755771636963
3	0.030130553990602493	0.2879081666469574	53	0.0002733312430791557	6.130360126495361


4	0.02610974945127964	0.020778097212314606	54	0.00012528635852504522	0.65610671043396
5	0.033010635524988174	0.31226906180381775	55	1.7217266758962069e-06	0.012183445505797863
6	0.0104494858533144	0.46299535036087036	56	4.478170376387425e-05	0.007874768227338791
7	0.015244405716657639	0.8511728048324585	57	0.0008844331023283303	0.0005371880251914263
8	0.0347551628947258	2.988057851791382	58	4.3432210077298805e-05	0.4318878948688507
9	0.06737023591995239	3.6387710571289062	59	6.217388727236539e-05	2.226496696472168
10	0.04552295804023743	4.034547328948975	60	0.12197347730398178	6.655332088470459
11	0.01012037880718708	3.177661657333374	61	0.00066091661574319	6.745500564575195
12	0.005894658155739307	3.7149665355682373	62	0.0007597299409098923	6.868861198425293
13	0.01004200242459774	2.4360475540161133	63	0.00044388126116245985	7.193081855773926
14	0.012006807141005993	3.505023241043091	64	0.0006670180591754615	7.340569019317627
15	0.0034727829042822123	4.801883697509766	65	0.000524028495419770	7.480850696563721
16	0.0030860528349876404	4.894974231719971	66	0.0004922249936498702	7.533853530883789
17	0.007208674680441618	3.50071382522583	67	0.0003795555967371911	7.639799118041992
18	0.001297667738981545	4.2584004402160645	68	0.0003163463552482426	8.026603698730469
19	0.02375160902738571	4.9730000495910645	69	0.00047755689593032	8.075765609741211
20	0.0025050409603863955	4.887083530426025	70	0.000378977885702624	8.051383018493652
21	0.0006188740371726453	0.6921811103820801	71	0.0002622345054987818	8.595427513122559
22	0.00023532958584837615	0.0003986225347034633	72	0.00038689016946591437	8.42041015625
23	0.0507756806910038	0.014116782695055008	73	0.00026363975484855473	8.243671417236328
24	7.433368591591716e-05	2.2276875972747803	74	0.00027741401572711766	8.171669006347656
25	0.09195014834403992	5.100812911987305	75	0.0702829509973526	8.50546646118164
26	0.004182533826678991	5.842406749725342	76	0.00018474162789061666	8.433987617492676
27	0.016741693019866943	7.258368968963623	77	0.00021110702073201537	8.01400375366211
28	0.003255910240113735	7.130593776702881	78	0.00023042417888063937	8.346099853515625


29	0.020350737497210503	6.613577365875244	79	0.00012125368812121451	7.620498180389404
30	0.0026462769601494074	6.50969123840332	80	0.0001418890169588849	8.030329704284668
31	0.0034328096080571413	6.755902290344238	81	0.0002293815341545269	8.174145698547363
32	0.0007909965352155268	5.331682205200195	82	0.00010832848056452349	8.840995788574219
33	0.001001332071609795	6.130487442016602	83	0.0001802610349841416	6.990046501159668
34	0.0024883649311959743	7.010079383850098	84	6.738227239111438e-05	7.423491954803467
35	0.0022047697566449642	6.4646992683410645	85	9.671755105955526e-05	7.959722995758057
36	0.0005545821622945368	5.107856273651123	86	0.0002011724136536941	8.179006576538086
37	0.006070754025131464	7.191289901733398	87	0.00016036634042393416	7.992383003234863
38	0.0004963439423590899	4.945166110992432	88	0.00017344091611448675	7.885098934173584
39	0.0007118553621694446	6.25137186050415	89	0.00011105283920187503	7.610461235046387
40	0.0005598401767201722	6.020808696746826	90	0.00275879236869514	8.571656227111816
41	0.013507477939128876	6.355642318725586	91	7.348756480496377e-05	7.771444320678711
42	0.00024395556829404086	5.49357795715332	92	0.0004482402582652867	8.002350807189941
43	0.0017429712461307645	7.3229475021362305	93	8.444989362033084e-05	8.586740493774414
44	0.00015546039503533393	4.242682933807373	94	0.0003593981673475355	0.01759181171655655
45	8.874354534782469e-05	5.051324367523193	95	2.7691252069139694e-11	7.074774657667149e-06
46	0.000597742444369942	6.1808953285217285	96	2.1056118839624105e-06	0.0
47	0.0005582966841757298	6.804682731628418	97	0.06361004710197449	1.995489927811289e-12
48	0.0004354499978944659	7.979450702667236	98	0.009731325320899487	8.20274042179122e-20
49	0.0011792771983891726	8.789504051208496	99	0.0008047168958000839	2.173821449279785




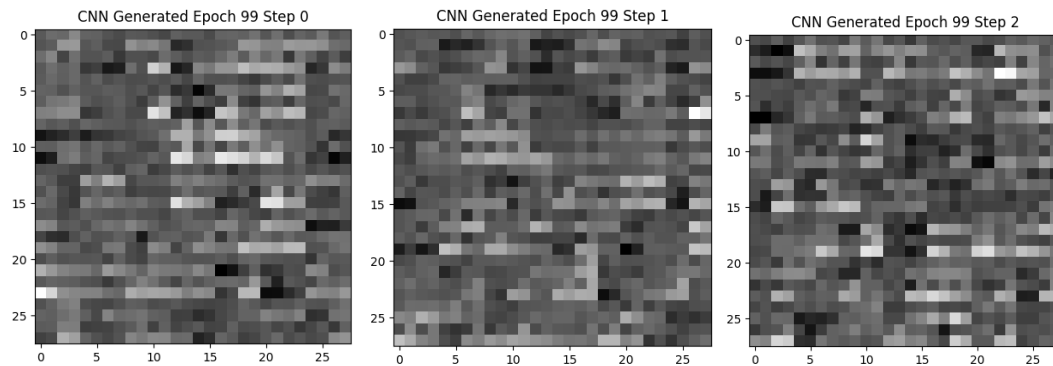
I plotted the loss to see what was happening, and it seems that my generator was not learning properly. I would have spent more time figuring out how to fix the issue, but I was short on time. There does seem to be a jump in loss for the discriminator (blue) when the loss for the CNN GAN model (orange) deeps down to its level, so I think there may be still competition between the two, but the discriminator seems to be just way better than the generator. The changes in the images below suggest that the generator is learning something but not quite what we wanted it to learn.

CNN Synthetic Images









Conclusion

The generator of the GAN model has learned that fashion items are lighter color and in the center area of the 28x28 square, but it still produced outputs with a lot of noise. I suppose the last epoch synthetic images pulled can be thought of as a shirt or sweater with a lot of noise. Generally, the generator seems to improve little by little with each epoch. The loss values fluctuate for the discriminator and GAN, which shows that the generator and the discriminator were in competition. I am not sure why I have GAN loss of 0 for epochs 1-9 for GAN in the beginning. It could be that the discriminator was really bad at spotting fake images and thought everything was real. The GAN model worked as expected. The CNN GAN model though, did not work as I expected and I think I either have architectural issues with the generator, or an overly good discriminator in comparison to the generator, or some kind of training issues with the generator. The generator was learning something, as it was producing different looking images per epoch, but it seems to be missing the point that the image is supposed to have a lighter colored object in the middle and the edges should be black.

References:

https://keras.io/guides/writing_a_training_loop_from_scratch/
<https://machinelearningmastery.com/practical-guide-to-gan-failure-modes/>